

NDN Packet Format Specification V0.1

NDN Project Team

November 1, 2013

Acknowledgment

This NDN packet format specification is adapted from the CCNx specification described in <http://www.ccnx.org/releases/latest/doc/technical/index.html> as of October 2013. However we have made a number of packet format changes based on our deepened understanding about NDN design through experimentation with NDN over the last three years, and based on the inputs from NDN research community at large. In particular, this specification adopted a TLV encoding format that has been championed by cisco NDN project team. We list the major protocol format changes from the CCNx specification at the end of each section.

1 Introduction

This version 0.1 specification aims to describe the NDN packet format only, a much narrower scope than a full NDN protocol specification. Our plan is to circulate and finalize the packet format first, then write down the full protocol specification.

In addition to this protocol specification draft, we are also in the process of putting out a set of technical memos that document our reasoning behind the design choices of important issues. The first few to come out will address the following issues:

- Packet fragmentation: end-to-end versus hop-by-hop;
- Understanding the tradeoffs of (not) handling Interest selectors;
- NDN Name discovery: why do we need it?
- NDN naming convention; and
- Scaling NDN routing.

In the rest of the document, we assume readers are familiar with how NDN/CCN works in general. For a description of the current CCNx protocol definition, please refer to <http://www.ccnx.org/releases/latest/doc/technical/CCNxProtocol.html>.

2 Type-Length-Value (TLV) Encoding

Each NDN packet is encoded in a Type-Length-Value (TLV) format as specified in Section 2.2. NDN Interest and Data packets are distinguished by the type value in the first and outmost TLV, TLV_0 .

An NDN packet is mainly a collection of TLVs inside TLV_0 . Some TLVs may contain sub-TLVs, and each sub-TLV may also be further nested. A guiding design principle is to keep the order of TLV_i s deterministic, and keep the level of nesting as small as possible to minimize both processing overhead and chances for errors.

Note that NDN packet format does not have a fixed packet header nor does it encode a protocol version number. Instead the design uses the TLV format to provide the flexibility of adding new types and phasing out old types as the protocol evolves over time. There is also no packet fragmentation support at network level. Whenever needed, NDN

packets may be fragmented and reassembled hop-by-hop.¹

2.1 Variable Size Encoding for type (T) and length (L)

(Both the text below and that in Sections 2.2 are adopted from an earlier packet specification draft by Mark Stapp)

To minimize the overhead during early deployment and to allow flexibility of future protocol extensions to meet unforeseeable needs, both type (T) and length (L) take a variable size format. For implementation simplicity, both type and length take the same encoding format.

We define a variable-length encoding for numbers in NDN as follows:

```
VAR-NUMBER := BYTE+
```

The first octet of the number either carries the actual numeric value, or signals that a multi-octet encoding is present, as defined below:

- if the first octet is < 253, the number is encoded in that octet;
- if the first octet == 253, the number is encoded in the following 2 octets, in net byte-order;
- if the first octet == 254, the number is encoded in the following 4 octets, in net byte-order;
- if the first octet == 255, the number is encoded in the following 8 octets, in net byte-order.

Each number must be encoded using the minimal number of octets possible to ensure unique encoding. This means that if a number is encoded in the X (= 2, 4, or 8) bytes following the first byte, its most significant byte must appear in the first X/2 bytes.

There are no special values or restrictions after the first VAR-NUMBER octet when the multiple byte encodings are used.

2.2 TLV Encoding

```
NDN-TLV := TLV-TYPE TLV-LENGTH TLV-VALUE?  
TLV-TYPE := VAR-NUMBER  
TLV-LENGTH := VAR-NUMBER  
TLV-VALUE := BYTE+
```

TLV-TYPE SHOULD be unique at all nested levels. The TLV Type number space and initial assignments will be specified elsewhere. NDN packet design will try best to keep the length of T staying with a single byte.

The TLV Length contains the number of bytes in the TLV including the length of the Type and Length fields. The Length field must not be zero. If a Type does not take any value, the Length field is 2.

This encoding offers a reasonable balance between compactness and flexibility. Most common, standardized Type codes will be allocated from a small-integer number-space, and these common Types will be able to use the compact, single-byte encoding.

2.3 Non Negative Integer Encoding

A number of TLV elements in NDN packet format take a non-negative integer as their value, with the following definition.

```
nonNegativeInteger ::= BYTE+
```

Length value of the TLV element MUST be either 3, 4, 6, or 10. Depending on the length value, a nonNegativeInteger is encoded as follows:

¹Today's IP networks provide point-to-point packet delivery and perform end-to-end fragmentation. An NDN network, on the other hand, may fetch requested data from any in-network storage, thus the notion of data flowing along an end-to-end path does not apply.

- if the length is 3 (i.e. the value length is 1 octet), the nonNegativeInteger is encoded in one octet;
- if the length is 4 (= value length is 2), the nonNegativeInteger is encoded in 2 octets, in net byte-order;
- if the length is 6 (= value length is 4), the nonNegativeInteger is encoded in 4 octets, in net byte-order;
- if the length is 10 (= value length is 8), the nonNegativeInteger is encoded in 8 octets, in net byte-order.

In the same as with VAR-NUMBER, each nonNegativeInteger must be encoded using the minimal number of octets possible to ensure unique encoding. This means that if a number is encoded using X (= 2, 4, or 8) bytes following the first byte, its most significant byte must appear in the first X/2 bytes.

The following shows a few examples of TLVs that has nonNegativeInteger as their value component in hexadecimal format (where TT represents TLV-TYPE, followed by the Length, then Value):

```
0      => TT0300
1      => TT0301
255    => TT03FF
256    => TT040100
65535  => TT04FFFF
65536  => TT0600010000
```

Changes from CCNx

- XML-based ccnb packet encoding is replaced by TLV encoding.

3 Name

An NDN Name is a hierarchical name for NDN content, which contains a sequence of name components.

3.1 NDN Name Format

We use a 2-level nested TLV to represent a name. The Type in the outer TLV indicates this is a Name. All inner TLVs have the same Type indicating that they each contain a name component. There is no restriction on the Value field in a name component and it may not contain any bytes.

```
Name ::= NAME-TYPE TLV-LENGTH NameComponent*
NameComponent ::= NAME-COMPONENT-TYPE TLV-LENGTH BYTE+
```

3.2 NDN URI Scheme

For textual representation, it is often convenient to use URI to represent NDN names. Please refer to RFC 3986 (URI Generic Syntax) for background.

- The scheme identifier is “ndn”.
- When producing a URI from an NDN Name, only the generic URI unreserved characters are left unescaped. These are the US-ASCII upper and lower case letters (A-Z, a-z), digits (0-9), and the four specials PLUS (+), PERIOD (.), UNDERSCORE (_), and HYPHEN (-). All other characters are escaped using either the percent-encoding method of the URI Generic Syntax or a “ndn” scheme specific hexadecimal string escape starting with the EQUALS (=) and an even number of characters from the set of hex digits. Once an EQUALS has been encountered in a component the hexadecimal encoding persists until the end of the component. The hex digits in these escaped encodings should always use upper-case letters, i.e., A-Z.
- To unambiguously represent name components that would collide with the use of . and .. for relative URIs, any component that consists solely of zero or more periods is encoded using three additional periods.
- The authority component (the part after the initial “//” in the familiar http and ftp URI schemes) is not relevant to NDN. It should not be present, and it is ignored if it is present.

3.3 Implicit Digest Component

The Name of every piece of content includes as its final component a derived digest that ultimately makes the name unique. This digest may occur in an Interest Name as an ordinary Component (the last one in the name). This final component in the name is never included explicitly in the Data packet when it is transmitted on the wire. It can be computed by any node based on the Data packet content.

The **implicit digest component** consists of the SHA-256 digest of the entire Data packet without the signature component. Having this digest as the last name component enables us to achieve the following two goals:

- Identify one specific Data packet and no other.
- Exclude a specific Data packet in an Interest (independent from whether it has a valid signature).

3.4 Canonical Order

In several contexts in NDN packet processing, it is useful to have a consistent ordering of names and name components. NDN names consist of a sequence of NameComponents, and each NameComponent is a sequence of zero or more 8-bit bytes. The ordering for components is such that:

- If a is shorter than b (i.e., has fewer bytes), then a comes before b .
- If a is the same length as, but different than, b , and i is the first place they differ, then a comes before b if and only if $a[i] < b[i]$ (treating $a[i]$ and $b[i]$ as unsigned 8-bit values). (This is just memcmp() ordering.)

For Names, the ordering is just based on the ordering of the first component where they differ. If one name is a proper prefix of the other, then it comes first.

Changes from CCNx

1. The name encoding is changed from binary XML to TLV format.
2. The discussions on naming conventions and the use of special markers inside NameComponents are removed from packet specification, and will be covered by a separate technical document ([1]).
3. Deprecated zero-length name component.

4 Interest Packet

Interest

```
Interest ::= INTEREST-TYPE TLV-LENGTH
           Name
           Selectors?
           Nonce
           Scope?
           InterestLifetime?
```

```
Selectors ::= SELECTORS-TYPE TLV-LENGTH
           MinSuffixComponents?
           MaxSuffixComponents?
           PublisherPublicKeyLocator?
           Exclude?
           ChildSelector?
           MustBeFresh?
```

Name and **Nonce** are the only two required elements in an Interest packet. Selectors are optional elements that further qualify Data that may match the Interest. They are used for discovering and selecting the Data that matches best to what the application wants. Selectors are placed right after the Name to facilitate implementations that may use continuous memory block of Name and Selectors TLVs together as the index for PIT lookup. By using a TLV to group all the Selectors, an implementation can easily skip them to find Nonce, which is used together with Name to identify looping Interests. If Selectors TLV is present in the Interest, it **MUST** contain at least one selector.

The two other optional elements, Scope and InterestLifetime, are referred to as *Guiders*. They affect Interest forwarding behavior, e.g., how far the Interest may be forwarded, and how long an Interest may be kept in the PIT. They are not grouped.

Name

The Name element in an Interest is synonymous with the term *prefix*. See Section 3 for details.

Nonce

Nonce ::= NONCE-TYPE TLV-LENGTH BYTE+

The Nonce carries a randomly-generated byte-string. The combination of Name and Nonce should uniquely identify an Interest packet. This is used to detect looping Interests².

MinSuffixComponents, MaxSuffixComponents

MinSuffixComponents ::= MIN-SUFFIX-COMPONENTS-TYPE TLV-LENGTH
nonNegativeInteger

MaxSuffixComponents ::= MAX-SUFFIX-COMPONENTS-TYPE TLV-LENGTH
nonNegativeInteger

When needed, **MinSuffixComponents** and **MaxSuffixComponents** allow a data consumer to indicate whether the Name in the Interest is the full name including the digest, or the full name except for the digest, or the content it is seeking has a known range of legitimate component counts. These two parameters refer to the number of name components beyond those in the prefix, and counting the implicit digest, that may occur in the matching Data. The default for **MinSuffixComponents** is 0 and for **MaxSuffixComponents** is effectively infinite, meaning that any Data whose name starts with the prefix is a match. Often only one of these will be needed to get the desired effect.

PublisherPublicKeyLocator

PublisherPublicKeyLocator ::= PUBLISHER-PUBLIC-KEY-TYPE TLV-LENGTH Name

This element specifies the name of the key which is used to sign the Data packet that the consumer is requesting. This is a way for the Interest to select answers from a particular publisher.

Exclude

Exclude ::= EXCLUDE-TYPE TLV-LENGTH Any? (NameComponent (Any)?)+
Any ::= ANY-TYPE TLV-LENGTH (=2)

²We plan to select a recommended length for Nonce in next revision.

The **Exclude** selectors allows requester to specify list and/or ranges of names components that **MUST NOT** appear as a continuation of the Name prefix in the responding Data packet to the Interest. For example, if Interest is expressed for `/ndn/edu` and Exclude specifies one name component `ucla`, then nor data producer nor conforming NDN routers are allowed to return any Data packet that has prefix `/ndn/edu/ucla`.

Exclude filter applies only to a name component of the Data packet name that is located at a position that numerically equals to the number of name components in the Interest packet, assuming 0 is the first name component.

The Components in the exclusion list **MUST** occur in strictly increasing order according to the canonical NDN name component ordering (Section 3), with optional leading, trailing, and interleaved “Any” components. The following defines processing of “Any” components:

- If none of the “Any” components are specified, the filter excludes only to the names specified in the Exclude list.
- If a leading “Any” component is specified, then the filter excludes all names that are smaller or equal (in NDN name component canonical ordering) to the first NameComponent in the Exclude list.
- If a trailing “Any” component is specified, then the filter excludes all names that are larger or equal (in NDN name component canonical ordering) to the last NameComponent in the Exclude list.
- If “Any” component is specified between two NameComponents in the list, then the filter excludes all names from the range from the right NameComponent to the left NameComponent, including both ends.

Exclude filter **MUST** not consist of a single “Any” component or one NameComponent with leading and trailing “Any” components.

ChildSelector

```
ChildSelector ::= CHILD-SELECTOR-TYPE TLV-LENGTH
                nonNegativeInteger
```

Often a given Interest can match more than one Data within a given content store. The **ChildSelector** provides a way of expressing a preference for which of these should be returned. If the value is 0, the leftmost child is preferred. If 1, the rightmost child is preferred. Here leftmost and rightmost refer to the least and greatest components according to the canonical NDN name component ordering (Section 3). This ordering is only done at the level of the name hierarchy one past the name prefix.

For example, assuming in the name hierarchy the component immediately after the name prefix is the version number, whose next level is the segment number, then setting ChildSelector to be 1 will retrieve the rightmost version number (i.e., the latest version) and the leftmost segment number (i.e., the first segment). However, this selection is only done with respect to a single content store, not globally. Additional rounds that exclude the earlier versions may be used to explore other content stores for newer versions. In this case, the use of ChildSelector does not change the multi-round outcome, but it decreases the number of rounds needed to converge to an answer.

MustBeFresh

```
MustBeFresh ::= MUST-BE-FRESH-TYPE TLV-LENGTH (=2)
```

This selector is encoded with Type and Length but no Value part. When it is absent from an Interest packet, the router can respond with a Data packet from its content store whose FreshnessPeriod is either still valid or expired. When it is present in an Interest packet, the router should not return Data packet from its content store whose FreshnessPeriod has expired.

The FreshnessPeriod carried in each Data packet (Section 5) is set by the original producer. It starts counting down when the Data packet arrives at a node. Consequently if a node is N hops away from the original producer, it may not consider the Data stale until $N \times$ FreshnessPeriod after the Data is produced.

Scope

```
Scope ::= SCOPE-TYPE TLV-LENGTH nonNegativeInteger
```

This value limits how far the Interest may propagate. Scope 0 prevents propagation beyond the local NDN daemon (even to other applications on the same host). Scope 1 limits propagation to the applications on the originating host. Scope 2 limits propagation to no further than the next node. Other values are not defined at this time, and will cause the Interest packet to be dropped.

Note that Scope is not a hop count - the value is not decremented as the Interest is forwarded.

InterestLifetime

`InterestLifetime ::= INTEREST-LIFETIME-TYPE TLV-LENGTH nonNegativeInteger`

InterestLifetime indicates the (approximate) time remaining before the Interest times out. The value is the number of milliseconds. The timeout is relative to the arrival time of the Interest at the current node.

Nodes that forward Interests may decrease the lifetime to account for the time spent in the node before forwarding, but are not required to do so. It is recommended that these adjustments be done only for relatively large delays (measured in seconds).

It is the application that sets the value for **InterestLifetime**. If the **InterestLifetime** element is omitted, a default value of 4 seconds is used (4000). The missing element may be added before forwarding.

Changes from CCNx

- Nonce is changed from optional to required.
- PublisherPublicKeyDigest is replaced by PublisherPublicKeyLocator.
- AnswerOriginKind is simplified from 4bits to a 1-bit MustBeFresh.
- FaceID has been removed.
- InterestLifetime changes the unit to the number of milliseconds.
- Removed Bloom Filter from Exclude.

5 Data Packet

Data

`Data ::= DATA-TLV TLV-LENGTH`
 Name
 MetaInfo
 Content
 Signature

The Data packet represents some arbitrary binary data (held in the Content element) together with its Name, some additional bits of information (MetaInfo), and a digital Signature of the other three elements. The Name is the first element since all NDN packet processing starts with the name. Signature is put at the end of the packet to ease the implementation because signature computation covers all the elements before Signature.

Name

See Section 3.

MetaInfo³

`MetaInfo ::= META-INFO-TYPE TLV-LENGTH`

³If both Content Type and Freshness Period are optional, one may consider MetaInfo itself should be optional. But would have all 4 parts of Data packet help simplify implementation? We leave this question to people who are more familiar with high speed implementations.

ContentType?
FreshnessPeriod?

Compared with CCNx, four fields are removed: PublisherPublicKeyDigest, ExtOpt, Timestamp, and FinalBlockID for the following reasons.

- PublisherPublicKeyDigest is supposed to be used in selecting data packets signed by a particular key. We replace PublisherPublicKeyDigest with KeyLocator, which is part of the Signature block (see Section 6), due to the following consideration. First, it requires data consumer to acquire a *valid* public key, as opposed to the key locator, before sending Interest out. Second, if a router is to verify the content objects, it must have other means to locate the keys first. Further, it may require publishers to maintain their public keys and certificates by their public key digests instead of names.
- ExtOpt was intended for extending XML-based ccnb format. Since we are now using TLV, ExtOpt is no longer needed.
- Timestamp and FinalBlockID can be useful meta information for applications, but do not need to be processed at the network layer. Therefore, if desired, applications should encode such meta information as part of the content.

ContentType

ContentType ::= CONTENT-TYPE-TYPE TLV-LENGTH nonNegativeInteger

Three ContentTypes are currently defined: default (=0), LINK (=1), and KEY (=2). The **default** type of content is a BLOB (=0), which is the actual data bits identified by the data name. The **LINK** type of content is another name which identifies the actual data content. The **KEY** type of content is a public key.

Compared with CCNx, three types, ENCR, GONE, and NACK are removed. ENCR means the content is encrypted, and since the network layer should not care whether content is encrypted or not, this type is not needed. GONE was a placeholder for implementing cache purging, however the research is yet to be carried out on how to accomplish this goal, if it is feasible to achieve, it is not included in this 0.1 version of NDN specification. NACK is used to signal a downstream node that the upstream node is unable to retrieve a matching data. Since the actual NACK mechanism is still under active investigation, we do not include it in this version of specification, but may add it back in a future version.

FreshnessPeriod

FreshnessPeriod ::= FRESHNESS-PERIOD-TLV TYPE-TLV nonNegativeInteger

The optional FreshnessPeriod indicates how long a node should wait after the arrival of this data before marking it as stale. The encoded value is number of milliseconds. Note that the stale data is still valid data; the expiration of FreshnessPeriod only means that the producer may have produced newer data.

Each content store associates every piece of Data with a staleness bit. The initial setting of this bit for newly-arrived content is “not stale”. If the Data carries FreshnessPeriod, then after the Data has been residing in the content store for FreshnessPeriod, it will be marked as stale. This is per object staleness and local to the NDN node. Another possible way to set the staleness bit of a local content is for a local client to send a command to the local NDN daemon.

If an Interest contains MustBeFresh TLV, a Data that has the staleness bit set is not eligible to be sent in response to that Interest. The effect is the same as if that stale Data did not exist (i.e., the Interest might be matched by some other Data in the store, or, failing that, get forwarded to other nodes). If an exact duplicate of a stale Data arrives, the effect is the same as if the stale Data had not been present. In particular, the Data in the store is no longer stale. As a practical matter, a stale Data should be ranked high on the list of things to discard from the store when a storage quota has been reached.

Content

Content ::= CONTENT-TYPE TLV-LENGTH BYTE*

Changes from CCNx

- The structure of Data packet is changed.
- SignedInfo is renamed to MetaInfo and its content is changed.
- PublisherPublicKeyDigest and ExtOpt are removed.
- Timestamp and FinalBlockID are removed.
- KeyLocator is moved to be inside the Signature block.
- Three content types, ENCR, GONE, and NACK are removed.
- FreshnessSeconds is renamed to FreshnessPeriod and is expressed in units of milliseconds.

6 Signature

```
Signature ::= SIGNATURE-TYPE TLV-LENGTH
    ( DigestSha256 |
      SignatureSha256WithRsa |
      SignatureSha256WithRsaAndMerkle |
      ...
    )
```

Signature TLV is a general container of signature, which contains an inner signature TLV. The type field of the inner signature TLV indicates the signing method of the signature, for example, DigestSha256 indicates that the integrity of Data is protected by a SHA-256 digest in DigestSha256; SignatureSha256WithRsa indicates that the integrity and provenance of Data is protected by a RSA signature over a SHA-256 digest; SignatureSha256WithRsaAndMerkle indicates that the integrity and provenance of Data is protected by a RSA signature over SHA-256-Merkle-Hash digest. If some other types of signatures are required, a new inner signature TLV will be defined.

For each inner signature TLV, the last embedded TLV must be a SignatureBits TLV, for example:

```
DigestSha256 ::= DIGEST-SHA256-TYPE TLV-LENGTH(=32) SignatureBits(=BYTE[32])
```

```
SignatureSha256WithRsa ::= SIGNATURE-SHA256-WITH-RSA-TYPE TLV-LENGTH
    KeyLocator
    SignatureBits(=BYTE[32])
```

```
SignatureSha256WithRsaAndMerkle ::= SIGNATURE-SHA256-WITH-RSA-AND-MERKLE-TYPE
    TLV-LENGTH
    KeyLocator
    Witness
    SignatureBits(=BYTE[32])
```

Which fields are covered by the SignatureBits TLV is up to the signing mechanism. A valid signing mechanism, however, must cover following TLVs: Name, MetaInfo (if present), and Content. Some signing mechanisms may also require the SignatureBits TLV to cover more TLVs. For example, SignatureSha256WithRsa requires the KeyLocator TLV to be signed, and SignatureSha256WithRsaAndMerkle requires both KeyLocator TLV and Witness TLV to be signed.

For inner signature TLVs that use public key cryptography, the first embedded TLV must be a KeyLocator TLV, e.g., as shown in SignatureSha256WithRsa and SignatureSha256WithRsaAndMerkle above.

```
KeyLocator ::= KEY-LOCATOR-TYPE TLV-LENGTH CertificateName |
    (other types of KeyLocators)
CertificateName ::= CERTIFICATE-NAME-TYPE TLV-LENGTH Name
```

A KeyLocator tells where to find the public key to verify this Data packet. For example, one can specify the name of the certificate of the public key (by CertificateName). Name conventions can be used to find the name of the key for a piece of content from the name of the Data packet.

All the other embedded TLVs in an inner signature TLV (such as Witness) are the signature-specific meta information and may vary from a signature type to another.

Changes from CCNx

- Signature is moved to the end of Data packet.
- KeyLocator is moved to be an inner signature block, making the inner signature block self-contained and self-sufficient.
- Signature type (or signing method information) is expressed by the type of inner signature TLV, rather than OID.
- Weak signatures, such as digest or checksum, are supported.
- ...

References

- [1] TBD. Ndn naming conventions.

DRAFT